

# churn\_prediction\_prepare\_data

January 23, 2023

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import OneHotEncoder, LabelEncoder
from datetime import date, timedelta
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
import imblearn
from imblearn.over_sampling import SMOTE
from sklearn.metrics import precision_score, recall_score, f1_score, \
    ↪ roc_auc_score
from sklearn.metrics import confusion_matrix
from xgboost import XGBClassifier
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import accuracy_score
from sklearn.ensemble import RandomForestClassifier
import datetime
from datetime import date, timedelta
```

## Data preparation and pre-processing

### Functions

```
[2]: ### Functions

def load_data(path):
    return (pd.read_csv(path))

def prepare_data(df_transactions, df_players):

    df_transactions.columns = ["x"]
    df_transactions = df_transactions["x"].str.split(';', expand = True)
    df_transactions.columns = ["player_id", "transaction_date", "product", \
    ↪ "transaction_type", "amount", "count"]

    #change types
    df_transactions['player_id'] = df_transactions['player_id'].astype('int32')
```

```

    df_transactions['transaction_date'] = df_transactions['transaction_date'].
↳astype('datetime64')
    df_transactions['product'] = df_transactions['product'].astype('string')
    df_transactions['transaction_type'] = df_transactions['transaction_type'].
↳astype('string')
    df_transactions['amount'] = df_transactions['amount'].astype('float32')
    df_transactions['count'] = df_transactions['count'].astype('int32')

    df_players.columns = ["player_id", "birth_date", "city",
↳"registration_date", "registration_hour", "is_opt_out",
↳"registration_terminal"]

    #change types
    df_players['birth_date'] = df_players['birth_date'].astype('datetime64')
    df_players['city'] = df_players['city'].astype('string')
    df_players['registration_date'] = df_players['registration_date'].
↳astype('datetime64')
    df_players['registration_hour'] = df_players['registration_hour'].
↳astype('float32')

    return df_transactions, df_players

def get_new_features(df):

    # Prepare for calculating new features later

    # Time since registration
    df['time_since_registration'] = (df['transaction_date'] -
↳df['registration_date']).dt.total_seconds()

    # Frequency and monetary value
    df['prev_amount'] = df.groupby('player_id')['amount'].shift(1)
    df['prev_amount'].fillna(df['amount'], inplace = True)

    df['frequency'] = (df.groupby('player_id')['transaction_date'].cumcount() +
↳1)

    df['monetary_value'] = df.groupby('player_id')['prev_amount'].cumsum()

    df.drop(['prev_amount'], axis = 1, inplace = True)

    return df

def clean_data(df):

```

```

len_before_clean = len(df)

# For some rows transaction_date is before registration_date which is not
↳ possible
df = df.drop(df[df['transaction_date'] < df['registration_date']].index)

# We have some null values (very small amount so best option is to delete)
df = df.dropna()

len_after_clean = len(df)

print("Removed ", (1 - len_after_clean / len_before_clean) * 100, "% of
↳ data.", sep = '')

return df

```

## Load dataset and prepare to calculate new features

```

[3]: df_transactions = load_data("dataset/zadatak-lite.csv")
df_players = load_data("dataset/igraci.csv")

df_transactions, df_players = prepare_data(df_transactions, df_players)

df = pd.merge(df_transactions, df_players, on = 'player_id', how = 'left')

df = clean_data(df)

df.sort_values(by = 'transaction_date', inplace = True)

```

Removed 5.001921881014326% of data.

```

[4]: df = get_new_features(df)

```

```

[5]: print(df['transaction_date'].min(), df['transaction_date'].max(), sep = '\n')

```

```

2022-04-01 00:00:00
2022-12-31 00:00:00

```

```

[6]: start_date = date(2022, 5, 1)
end_date = date(2022, 11, 30)
date_range = [start_date + timedelta(days=x) for x in range((end_date -
↳ start_date).days + 1)]

new_df = pd.DataFrame(columns = ['player_id', 'date', 'bo_count',
↳ 'casino_count', 'pp_count', 'sport_count', 'vb_count',
    'vdr_count', 'frequency', 'monetary_value', 'profit', 'deposit',
    'last_active', 'player_age', 'time_since_registration', 'is_opt_out',

```

```

        'churn']))

for d in date_range:
    #print(d)
    d = pd.to_datetime(d)

    last_30_days_df = df[(df['transaction_date'] >= d - timedelta(days = 30)) &
↳(df['transaction_date'] < d)]

    # Count by product
    # Group by player_id, product
    product_count_by_player = last_30_days_df.
↳groupby(['player_id', 'product'])['count'].sum().reset_index()
    product_count_by_player = product_count_by_player.pivot(index =
↳'player_id', columns = 'product', values = 'count')
    # Fillna with 0
    product_count_by_player.fillna(0, inplace=True)
    product_count_by_player.rename(columns = {'Sport': 'sport_count', 'Casino':
↳'casino_count',
                                                    'PaymentProvider':
↳'pp_count', 'BusinessOwner': 'bo_count',
                                                    'VirtualBingo': 'vb_count',
↳'VirtualDogRace': 'vdr_count'},
                                inplace = True)

    # Frequency
    frequency_by_player = last_30_days_df.groupby(['player_id'])['frequency'].
↳max() - last_30_days_df.groupby(['player_id'])['frequency'].min()

    frequency_by_player = frequency_by_player.to_frame()
    frequency_by_player = frequency_by_player.rename(columns={0: "frequency"})

    # Monetary value
    monetary_by_player = last_30_days_df.
↳groupby(['player_id'])['monetary_value'].max() - last_30_days_df.
↳groupby(['player_id'])['monetary_value'].min()

    monetary_by_player = monetary_by_player.to_frame()
    monetary_by_player = monetary_by_player.rename(columns={0:
↳"monetary_value"})

    # Profit, Deposit
    transaction_amount_by_player = last_30_days_df.groupby(['player_id',
↳'transaction_type'])['amount'].sum().reset_index()

```

```

transaction_amount_by_player = transaction_amount_by_player.pivot(index =
↳ 'player_id', columns = 'transaction_type', values = 'amount')
transaction_amount_by_player.fillna(0, inplace = True)
transaction_amount_by_player.rename(columns = {'Bonus': 'bonus',
↳ 'TicketWin': 'ticketwin', 'TicketPayin': 'payin', 'Deposit':
↳ 'deposit', 'Withdrawal': 'withdrawal'}, inplace = True)
profit = transaction_amount_by_player['ticketwin'] +
↳ transaction_amount_by_player['bonus'] - transaction_amount_by_player['payin']
transaction_amount_by_player.insert(0, 'profit', profit)
profit_by_player = transaction_amount_by_player.drop(columns = ['bonus',
↳ 'payin',
                                'ticketwin', 'withdrawal', 'DepositCancel',
↳ 'TicketPayinCancel', 'TicketWinCancel'])

# Last active
last_active_by_player = (d - last_30_days_df.
↳ groupby(['player_id'])['transaction_date'].max())

last_active_by_player = last_active_by_player.to_frame()
last_active_by_player = last_active_by_player.
↳ rename(columns={'transaction_date': "last_active"})

# Player age
player_age = d - last_30_days_df.groupby(['player_id'])['birth_date'].max()

player_age = player_age.to_frame()
player_age = player_age.rename(columns={'birth_date': "player_age"})

# Time since registration
time_since_registration = d - last_30_days_df.
↳ groupby(['player_id'])['registration_date'].max()

time_since_registration = time_since_registration.to_frame()
time_since_registration = time_since_registration.
↳ rename(columns={'registration_date': "time_since_registration"})

# Is opt out
is_opt_out = last_30_days_df.groupby(['player_id'])['is_opt_out'].max()

is_opt_out = is_opt_out.to_frame()
is_opt_out = is_opt_out.rename(columns={0: "is_opt_out"})

# Merge all features in one dataframe
concat_df = pd.concat([product_count_by_player, frequency_by_player,
↳ monetary_by_player,
                                profit_by_player, last_active_by_player, player_age,

```

```

time_since_registration, is_opt_out], axis=1)

# Fix for player_id column
concat_df['player_id'] = concat_df.index

# Churn
concat_df['churn'] = concat_df['player_id'].isin(df[(df['transaction_date'] <=
→ d) &
(df['transaction_date'] < d + timedelta(days = 30))]['player_id']).astype(int)

concat_df['churn'] = concat_df['churn'].replace({0:1, 1:0})

# Date
concat_df['date'] = d

# Append to complete dataframe

new_df['frequency'] = new_df['frequency'].astype('int')
new_df['is_opt_out'] = new_df['is_opt_out'].astype('int')
new_df['player_id'] = new_df['player_id'].astype('int')
new_df['churn'] = new_df['churn'].astype('int')

new_df = pd.concat([new_df, concat_df], ignore_index = True)

new_df = new_df.drop('WithdrawalCancel', axis = 1)
new_df.head()

```

```

[6]:
  player_id      date  bo_count  casino_count  pp_count  sport_count  \
0         2  2022-05-01        1.0           0.0         1.0         13.0
1         4  2022-05-01       11.0        16992.0        79.0         25.0
2         5  2022-05-01         2.0        19796.0        43.0          0.0
3         8  2022-05-01         0.0           0.0         1.0          8.0
4        11  2022-05-01         1.0           0.0         0.0          0.0

  vb_count  vdr_count  frequency  monetary_value      profit  deposit  \
0         0.0         0.0         2         38.009998    -1.809998    1.810000
1       6153.0         0.0       167        2495.419922   -258.329956   257.429993
2          0.0         0.0        28        7263.290039  -411.770020   592.309998
3          0.0         0.0         7          4.910000   -1.630000    0.900000
4          0.0         0.0         0          0.000000    4.520000    0.000000

  last_active  player_age  time_since_registration  is_opt_out  churn
0      22 days  10168 days           2777 days             0       0
1       2 days  18747 days           2777 days             0       0
2       6 days  15087 days           2777 days             0       0

```

3	1 days	18619 days	2776 days	0	0
4	23 days	9963 days	2776 days	0	1

### Save csv file

```
[7]: # Save dataframe ready for model
tmp_df = new_df
tmp_df.to_csv('dataset/new_feature_dataset.csv')
```

# churn\_prediction\_model

January 23, 2023

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.preprocessing import OneHotEncoder, LabelEncoder
from datetime import date, timedelta
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
import imblearn
from imblearn.over_sampling import SMOTE
from sklearn.metrics import precision_score, recall_score, f1_score, \
    roc_auc_score
from sklearn.metrics import confusion_matrix
from xgboost import XGBClassifier
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import accuracy_score
from sklearn.ensemble import RandomForestClassifier
import datetime
from datetime import date, timedelta
```

## Load dataset

```
[14]: df = pd.read_csv('dataset/new_feature_dataset.csv')
df.drop(['Unnamed: 0'], axis=1, inplace=True)
df['last_active'] = df['last_active'].apply(lambda x: int(x.split(" ")[0]))
df['player_age'] = df['player_age'].apply(lambda x: int(x.split(" ")[0]))
df['time_since_registration'] = df['time_since_registration'].apply(lambda x: \
    int(x.split(" ")[0]))
df.head()
```

```
[14]:
```

	player_id	date	bo_count	casino_count	pp_count	sport_count	\
0	2	2022-05-01	1.0	0.0	1.0	13.0	
1	4	2022-05-01	11.0	16992.0	79.0	25.0	
2	5	2022-05-01	2.0	19796.0	43.0	0.0	
3	8	2022-05-01	0.0	0.0	1.0	8.0	
4	11	2022-05-01	1.0	0.0	0.0	0.0	

  

	vb_count	vdr_count	frequency	monetary_value	profit	deposit	\
--	----------	-----------	-----------	----------------	--------	---------	---



0	0.0	0.0	2	38.01	-1.809998	1.81
1	6153.0	0.0	167	2495.42	-258.329960	257.43
2	0.0	0.0	28	7263.29	-411.770020	592.31
3	0.0	0.0	7	4.91	-1.630000	0.90
4	0.0	0.0	0	0.00	4.520000	0.00

	last_active	player_age	time_since_registration	is_opt_out	churn
0	22	10168	2777	0	0
1	2	18747	2777	0	0
2	6	15087	2777	0	0
3	1	18619	2776	0	0
4	23	9963	2776	0	1

Split the data into training and test sets

```
[3]: def split_df_by_date(df, date_col, date):
      df_before = df[df[date_col] < date]
      df_after = df[df[date_col] >= date]
      return df_before, df_after
```

```
[4]: df_train, df_test = split_df_by_date(df, 'date', '2022-10-30')
```

```
[5]: #X_train = df_train[['bo_count', 'casino_count', 'pp_count', 'sport_count',
      ↪ 'vb_count', 'vdr_count',
      #
      ↪ 'frequency', 'monetary_value', 'profit', 'deposit',
      ↪ 'last_active', 'player_age',
      #
      ↪ 'time_since_registration', 'is_opt_out']]
#X_test = df_test[['bo_count', 'casino_count', 'pp_count', 'sport_count',
      ↪ 'vb_count', 'vdr_count',
      #
      ↪ 'frequency', 'monetary_value', 'profit', 'deposit',
      ↪ 'last_active', 'player_age',
      #
      ↪ 'time_since_registration', 'is_opt_out']]

X_train = df_train[['frequency', 'monetary_value', 'last_active']]
X_test = df_test[['frequency', 'monetary_value', 'last_active']]

y_train = df_train[['churn']].values.ravel()
y_test = df_test[['churn']].values.ravel()

smote = SMOTE(sampling_strategy = "minority")
```

Logistic Regression

```
[6]: clf = LogisticRegression(penalty = None, random_state = 42, max_iter = 10000)
      clf.fit(X_train, y_train)
```

```
[6]: LogisticRegression(max_iter=10000, penalty=None, random_state=42)
```

## Optimal threshold

Ovdje smo optimizirali tako da zahtijevamo da barem 85% pravih churnera naš model detektira kao churnere, a u isto vrijeme minimiziramo našu metriku koja nam govori koliko smo ukupno igrača detektirali kao churnere. Recimo da šaljemo bonuse svim igračima koje model detektira kao churnere. Ovisno o tome koliko profita nam donosi ispravno detektiranje churnera i koliki je gubitak ako non-churnera detektiramo kao churnera ili churnera kao non-churnera mijenjali bismo fiksni postotak koji je ovdje 85%.

```
[7]: #thresholds = np.arange(0, 1.05, 0.01)
thresholds = np.arange(0.14, 0.16, 0.0001)

metrics = []

for threshold in thresholds:

    y_pred = (clf.predict_proba(X_test)[: ,1] >= threshold)
    tn, fp, fn, tp = confusion_matrix(y_test, y_pred).ravel()

    custom_metric = (tp + fp) / (tp + fp + tn + fn)
    rec = tp / (tp + fn)

    if(rec > 0.85):
        metrics.append(custom_metric)
    else:
        metrics.append(2.0)

best_threshold = thresholds[np.argmin(metrics)]
best_custom_metric = np.min(metrics)

y_pred = (clf.predict_proba(X_test)[: ,1] >= best_threshold)

print("Best threshold:", best_threshold)
print("Best custom metric:", best_custom_metric)
```

Best threshold: 0.14459999999999995

Best custom metric: 0.3305623744699844

## Optimize company profit

```
[8]: # Send bonus to churner
profit_tp = 5.0
# Send bonus to non_churner
profit_fp = -1.0
# Did not send bonus to churner
profit_fn = -5.0

#thresholds = np.arange(0, 1.05, 0.01)
thresholds = np.arange(0.11, 0.13, 0.0001)
```

```

profits = []

for threshold in thresholds:

    y_pred = (clf.predict_proba(X_test)[: ,1] >= threshold)
    tn, fp, fn, tp = confusion_matrix(y_test, y_pred).ravel()

    profit = profit_tp * tp + profit_fp * fp + profit_fn * fn

    profits.append(profit)

best_threshold = thresholds[np.argmax(profits)]
best_profit = np.max(profits)

y_pred = (clf.predict_proba(X_test)[: ,1] >= best_threshold)

print("Best threshold:", best_threshold)
print("Best profits:", best_profit)

```

Best threshold: 0.125400000000000046

Best profits: 55781.0

### Calculate important metrics

```

[9]: # calculate accuracy
accuracy = clf.score(X_test, y_test)
print("Accuracy: {:.2f}%".format(accuracy*100))

# calculate precision
precision = precision_score(y_test, y_pred)
print("Precision: ", precision)

# calculate recall
recall = recall_score(y_test, y_pred)
print("Recall: ", recall)

# calculate f1-score
f1 = f1_score(y_test, y_pred)
print("F1-Score: ", f1)

# calculate AUC-ROC
auc_roc = roc_auc_score(y_test, y_pred)
print("AUC-ROC: ", auc_roc)

```

Accuracy: 88.04%

Precision: 0.2912517918716826

Recall: 0.8754257765873824

F1-Score: 0.43708618897763685

AUC-ROC: 0.7927765680528369

### Confusion matrix

```
[10]: cm = confusion_matrix(y_test, y_pred)

# plot the confusion matrix
sns.heatmap(cm, annot=True, fmt="d", annot_kws={"size": 30})

plt.xticks(fontsize = 30)
plt.yticks(fontsize = 30)

plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()
```



### Correlation matrix

```
[11]: # Adjust the size of the figure
plt.figure(figsize = (20, 10))

# Create the heatmap with larger annotations
```

```
sns.heatmap(df.corr(), annot = True, fmt = ".2f", cmap = "YlGnBu", linewidths=.
↪5, annot_kws = {"size": 20})

# Increase the font size of the labels on the x and y axes
plt.xticks(fontsize = 20)
plt.yticks(fontsize = 20)

# Show the plot
plt.show()
```

/tmp/ipykernel\_40049/2151628175.py:5: FutureWarning: The default value of numeric\_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric\_only to silence this warning.

```
sns.heatmap(df.corr(), annot = True, fmt = ".2f", cmap = "YlGnBu",
linewidths=.5, annot_kws = {"size": 20})
```

